

COSC264 · INTRODUCTION TO COMPUTER NETWORKS AND THE INTERNET · TERM 4 · MODULE 6

# Application Protocols: the Web, Email and DNS

1 · Network applications

2 · The Web and HTTP

3 · Email

4 · DNS

Module 6 – Application Protocols

[cosc264.up.railway.app](https://cosc264.up.railway.app)

9 Oct 2026 · 1 / 50

OUTLINE

# Module 6 — roadmap

**General concepts about network applications**

**One network application: the Web**

**HTTP: the protocol that supports the Web**

**Email**

**DNS: the Domain Name System**

MODULE 6

# Network Applications

---

What an application needs from the layers beneath it

1 • Network applications

2 • The Web and HTTP

3 • Email

4 • DNS

# What is a network application?

**Programs that run on different end systems** and communicate with each other over the network.

## Web browsing

Google Chrome, Mozilla Firefox,  
Microsoft Edge

## Email

Gmail, Outlook

## Instant messaging

Slack, WhatsApp, Facebook  
Messenger

## Streaming services

YouTube, Netflix

## Social media

Instagram, TikTok

## AI assistants

Claude, ChatGPT, Gemini

# Application-layer protocols

An **application-layer protocol** is one piece of a network application — but a big piece.

It defines how the application's processes, running on different end systems, **pass messages to each other**: which messages exist, their format, and when and how to respond.

This is what lets programs on different machines work together — e.g. any web browser can fetch a page from any web server, whoever wrote each program.

| Network application | Application-layer protocol |
|---------------------|----------------------------|
| Web application     | HTTP                       |
| Email               | SMTP                       |

# Choosing a transport protocol

- Application-layer protocols are controlled by the **application developer**; transport-layer protocols, however, are built into the **OS**.
- As a developer you can only choose between the transport protocols the OS provides — so pick the one that best fits what your application-layer protocol needs.

| Property               | Example applications                                | Transport protocol                                      |
|------------------------|-----------------------------------------------------|---------------------------------------------------------|
| Reliable data transfer | Email, IM, file transfer, financial apps            | <b>TCP</b> preferred<br>UDP unsuitable                  |
| Low delay              | Internet telephony, conferencing, multiplayer games | <b>UDP</b> preferred<br>TCP retransmissions add latency |

MODULE 6

# The Web and HTTP

---

The application-layer protocol behind every page you load

1 · Network applications

2 · **The Web and HTTP**

3 · Email

4 · DNS

# The World Wide Web

## THE WEB

The **World Wide Web** (aka Web) is a **network application**: a global collection of documents and other resources shared over the Internet — like a large file storage system.

Each object on the Web (e.g., an image, an HTML file) is addressable by a **URL**:

```
http://www.someschool.edu/someDept/pic.gif  
└─ host name ─┐ └─ path name ─┘
```

## HTTP

# HTTP: the protocol behind the Web

### HTTP

The **Hypertext Transfer Protocol** is the **application-layer protocol** that supports the Web — it defines how clients request objects by URL and how servers respond.

HTTP follows a **request–response** pattern: the client sends an **HTTP request** for an object, and the server answers with an **HTTP response** carrying it.



## HTTP

# How HTTP works

HTTP uses **TCP** as its transport protocol — so before any HTTP message is sent, a TCP connection must be established first.



### STATE LESS

HTTP is **stateless** — the server maintains no information about past client requests. Each request is handled independently.

# HTTP connections

## Non-persistent HTTP

- At most **one object** is sent over a TCP connection — connection is then closed
- HTTP/1.0 uses non-persistent HTTP

## Persistent HTTP

- **Multiple objects** can be sent over a single TCP connection
- HTTP/1.1, HTTP/2, and HTTP/3 use persistent HTTP

| Protocol | Published | Connection type   | Status   |
|----------|-----------|-------------------|----------|
| HTTP/1.0 | 1996      | Non-persistent    | Legacy   |
| HTTP/1.1 | 1997      | Persistent        | Standard |
| HTTP/2   | 2015      | Persistent        | Standard |
| HTTP/3   | 2022      | Persistent (QUIC) | Standard |

# Non-persistent HTTP in action

Suppose a client wants to retrieve a file from: `http://www.someSchool.edu/someDepartment/home.index`

1. Client establishes a **TCP connection** with the server
2. Client sends an **HTTP GET** request — path: `/someDepartment/home.index`
3. Server retrieves the HTML file, wraps it in an **HTTP response**, and sends it back
4. Server signals TCP to **close the connection**



# Response time: non-persistent HTTP

## RESPONSE TIME

The time to retrieve a single object using HTTP — from the client's request until the object is received

**Q:** What is the response time for non-persistent HTTP?

**A:** 2 RTTs + file transmission time

## RTT

Recall: **round-trip time** — the time for a small packet to travel from client to server and back

**First RTT:** the TCP handshake — client sends SYN, server replies SYN-ACK

**Second RTT:** client sends HTTP GET, server sends back the response

With many referenced objects, this cost multiplies — a significant burden on the server



## Persistent HTTP in action

Client and server establish a TCP connection and exchange HTTP messages

Server **keeps the TCP connection open** after sending the first HTTP response

Subsequent HTTP messages between the same client and server use the **same TCP connection** — no new handshake needed

### Response time on one persistent connection:

1 RTT for the initial TCP handshake, then 1 RTT for each HTTP request.

So for  $n$  requests: **total**  $\approx (1 + n)$  RTT (plus transmission time).

Amortised over the  $n$  objects:  $(1 + n)/n = (1/n + 1)$  RTT. As  $n$  becomes large,  $1/n \rightarrow 0$ , so this approaches **~1 RTT per object**.

Can we do even better than 1 RTT per object?



HTTP

## Improving HTTP efficiency further with pipelining

**WITHOUT  
PIPELINING**

Client issues a new request **only after** receiving the previous response —  
**1 RTT per file**



# Improving HTTP efficiency further with pipelining

## WITHOUT PIPELINING

Client issues a new request **only after** receiving the previous response —  
**1 RTT per file**

## WITH PIPELINING

Client sends requests **as soon as it needs to retrieve an object**, without waiting for responses — e.g. in this example we pipeline 3 HTTP requests all together

**Benefit:** pipelining 3 requests → response time per file  $\approx 1 \text{ RTT} / 3$

- Don't confuse this with TCP-level pipelining.
- **HTTP pipelining** operates at the **application layer**: the client sends several **HTTP requests** without waiting for each response.
- **TCP pipelining** operates at the **transport layer**: the sender keeps several **segments** in flight within its window. TCP cuts each HTTP message into segments, so these may be pieces of one HTTP request, or of several.



# HTTP request message format

Method — what action to perform

| Method        | Action                          |
|---------------|---------------------------------|
| <b>GET</b>    | Retrieve a resource             |
| HEAD          | Retrieve headers only (no body) |
| <b>POST</b>   | Submit data to a resource       |
| PUT           | Store / replace a resource      |
| <b>DELETE</b> | Remove a resource               |

Request line

method URL version cr lf

Header lines

field name: value cr lf  
field name: value cr lf

Blank line

cr lf

Entity body

*entity body*

# HTTP request message format

**URL:** path to the requested resource — e.g.

`/someDept/pic.gif`

**HTTP Version:** identifies the protocol in use — HTTP/1.0 or HTTP/1.1

**Request line**

method

URL

version

cr lf

**Header lines**

field name:

value

cr lf

field name:

value

cr lf

**Blank line**

cr lf

**Entity body**

*entity body*

# HTTP request message format

## Header lines

- Zero or more `field-name: value` pairs, one per line
- Common examples:

```
Host: www.example.com
```

Mandatory in HTTP/1.1 — identifies the target server

```
Content-Type: text/html; charset=UTF-8
```

Format of the message body

```
Connection: close
```

Close the TCP connection after this response

```
Accept-Language: da, en-gb;q=0.8,  
en;q=0.7
```

Preferred languages, weighted by q (default 1): Danish → British English → any English

## Request line

method URL version cr lf

## Header lines

field name: value cr lf

field name: value cr lf

## Blank line

cr lf

## Entity body

*entity body*

HTTP

## HTTP request: an example

```
GET /somedir/page.html HTTP/1.1
```

```
Host: www.someschool.edu
```

```
User-agent: Mozilla/4.0
```

```
Connection: close
```

```
Accept-language: fr
```

*(empty entity body)*

**Method:** GET

**URL:** /somedir/page.html

**HTTP version:** HTTP/1.1

**Header field:** Host: www.someschool.edu

**Header field:** User-agent: Mozilla/4.0

**Header field:** Connection: close

**Header field:** Accept-language: fr

**Body:** (empty)

# HTTP response message format

## Version

The HTTP version the server is using — e.g. HTTP/1.1, just as in the request line.

Status line

version

status code

phrase

cr lf

Header lines

field name:

value

cr lf

field name:

value

cr lf

Blank line

cr lf

Entity body

*entity body*

# HTTP response message format

## Status code & phrase

- **Status code:** 3-digit number indicating the result of the request
- **Phrase:** human-readable text accompanying the code — e.g. OK, Not Found

## Status line

version — status code — phrase — cr lf

## Header lines

field name: — value — cr lf

field name: — value — cr lf

## Blank line

cr lf

## Entity body

*entity body*

# HTTP response message format


## Common examples

| Code | Phrase                     |
|------|----------------------------|
| 200  | OK                         |
| 301  | Moved Permanently          |
| 400  | Bad Request                |
| 404  | Not Found                  |
| 505  | HTTP Version Not Supported |


### Status line




### Header lines



### Blank line



### Entity body



# HTTP response message format

## Header lines, blank line, entity body

Same structure as in the request: `field-name: value`  
header lines, a blank line (a lone CRLF), then the body —  
here, the **requested object** itself.



# HTTP response: an example

```
HTTP/1.1 200 OK
Connection: close
Date: Thu, 06 Aug 1998 12:00:15 GMT
Server: Apache/1.3.0 (Unix)
Last-Modified: Mon, 22 Jun 1998 ...
Content-Length: 6821
Content-Type: text/html
```

```
<!DOCTYPE html>
<html><head><title>Example</title></head>
<body><h1>Hello!</h1>
<script>console.log("hi");</script>
</body></html>
```

**Version:** HTTP/1.1

**Status code:** 200 — success

**Phrase:** OK

**Header:** Connection: close — close after this response

**Header:** Date — when the response was sent

**Header:** Server — the server software

**Header:** Last-Modified — when the object last changed

**Header:** Content-Length — body size in bytes

**Header:** Content-Type — the body is HTML

**Body:** the requested HTML file (with embedded JS)

# Cookies

HTTP is **stateless** — every request is independent; the server retains no memory of past requests.

So the server does not remember your earlier conversations with it. Yet many applications *need* that memory: when you refresh an online shop's page, the items in your cart should not just disappear. **How does the server know what is in your cart?**

The answer: **cookies** — here's how it works.

1. A **Set-Cookie** header line in the **HTTP response** — the server sends the client a unique ID on first contact
2. A **Cookie** header line in subsequent **HTTP requests** — the client echoes the ID back on every future request
3. A **cookie file** stored on the client, managed by the browser — persists across sessions
4. A **back-end database** on the server, keyed by the cookie ID — stores per-user state

HTTP

## Cookies: an example [RFC 6265]



The client sends an HTTP request to the host

**Set-Cookie: id=1678** — Server assigns a unique ID and sends it to the client

The cookie ID is stored as an entry in the back-end database, and in the client's browser

# Cookies: an example [RFC 6265]



**Cookie: id=1678** — Client attaches the ID on every future request to the same server

Server uses the ID to look up the user's state in its database

Server can then generate a **tailored response** — e.g. a personalised page, a pre-filled shopping cart — for the client's request

HTTP

## One problem with the Web

Picture a university: thousands of users fetch the same popular pages and images, and every request travels across the Internet to the **origin server** that hosts the object.

Every one of those requests lands on the same origin server, through the same congested links — enough of them can **overwhelm** it.

**How can we avoid this?**

# Web caching

A **web cache** (also called a **proxy server**) is a network entity that satisfies HTTP requests on behalf of an origin server — storing copies of recently requested objects locally.



## Acts as a server

Receives HTTP requests from clients and, if it holds a copy of the object, returns it in the response — just like an origin server would

## Acts as a client

When the object is not cached, it issues its own HTTP request to the origin server, stores the copy, then answers the client

Requests answered by the cache get a **quicker response**, and caching **reduces how often** requests have to go to the origin server.

# Conditional GET



A **conditional GET** is an ordinary GET with an `If-modified-since` header line. A proxy uses it to check whether its cached copy is still up to date, e.g.:

```
GET /index.html HTTP/1.1
```

```
Host: www.example.com
```

```
If-modified-since: Mon, 10 Aug 2026 09:00:00  
GMT
```

If the object is **unchanged**, server replies:

```
304 Not Modified
```

Object **changed**: server replies 200 OK with the updated object in the body

MODULE 6

# Email

---

Electronic Mail


1 · Network applications

2 · The Web and HTTP

**3 · Email**

4 · DNS

# An overview of the email system



**User agent (UA)** — the app used to read, compose and send email (e.g. Outlook, Gmail, Apple Mail)

**Mail server** — a dedicated host that handles email on behalf of its users. It stores a **mailbox** per recipient for incoming mails and a **send queue** of outgoing mails

**SMTP** — Simple Mail Transfer Protocol; the sender's UA uses it to send a message to its own mail server

The sender's mail server then uses **SMTP** to transfer the message to the recipient's mail server across the internet

**Mail access protocol (MAP)** — used by UAs to retrieve messages from their mail server

## SMTP in action

SMTP has a **client side** and a **server side**, and always runs over **TCP**.



① Alice uses her UA to **compose** the message

② Alice's UA (**SMTP client**) **sends** the message to her mail server (**SMTP server**)

③ The **SMTP client** on Alice's server opens a TCP connection to Bob's mail server and **transfers** Alice's message over it

④ The **SMTP server** on Bob's server receives the message and places it in Bob's **mailbox**

## Retrieving email from the server



The recipient's UA retrieves email from its mail server either with a **mail access protocol (MAP)**, e.g.

- **POP3** — Post Office Protocol version 3 [RFC 1939]
- **IMAP** — Internet Message Access Protocol [RFC 9051]

or through **web-based email**, where the UA is a web browser that uses **HTTP** to reach the mail server

# How POP3 works



IMAP

# Internet Message Access Protocol

IMAP addresses POP3's limitations:

IMAP **keeps per-message state on the server**, e.g. read or unread, flagged, so every device sees the same status

It supports **folders on the server**: create folders and move messages into them, visible from any device

# Accessing email via a browser

## Web-based email is accessed through a web browser

The mailbox stays on the server; reading an email is like browsing the Web — the browser fetches each message as a web page

## Which application-layer protocol does the browser use to talk to the mail server?

**HTTP**, the same protocol behind every other web page

## Web-based email services

Examples: **Gmail**, **Outlook** (web)

## Comparing POP3, IMAP and web-based email

|                         | POP3                                                                  | IMAP                                                | Web-based email                                 |
|-------------------------|-----------------------------------------------------------------------|-----------------------------------------------------|-------------------------------------------------|
| <b>Accessed with</b>    | A mail app                                                            | A mail app                                          | A web browser, over <b>HTTP</b>                 |
| <b>Messages</b>         | Downloaded to the device; deleted from the server, or optionally kept | Kept on the server                                  | Kept on the server                              |
| <b>Folders</b>          | None on the server; a single mailbox                                  | Created and kept on the server                      | Created and kept on the server                  |
| <b>Multiple devices</b> | Not synced: read status and folders exist only on each device         | Read status and folders stay in sync across devices | Any device with a browser sees the same mailbox |

# Structure of an email message

An email message consists of **header fields**, followed (optionally) by a **body**

**Header fields** describe the message, e.g. **From:** and **Subject:**

The **body** carries the message content itself

The whole message is **ASCII text**, which encodes English letters, digits and punctuation

**Problem:** ASCII alone cannot carry images, attachments, or non-ASCII characters (e.g. Chinese, Arabic)

Date: Mon, 10 Aug 2026 09:00 GMT

From: alice@crepes.fr

To: bob@hamburger.edu

Subject: Picture of yummy crepe.

Cc: charlie@example.com

**Header fields**

---

Hello Bob, here is the picture  
you asked for. Enjoy!

**Body**

## MIME

# Multipurpose Internet Mail Extensions

**MIME** extends the message format with **additional header fields**

These fields declare the content's **type** and how it is **encoded**

so the receiver knows how to **decode the body**

e.g. **base64** encoding converts binary data into printable ASCII text, which can be used for image attachments

From: alice@crepes.fr

To: bob@hamburger.edu

Subject: Picture of yummy crepe.

**MIME-Version: 1.0**

**Content-Transfer-Encoding: base64**

**Content-Type: image/jpeg**

(base64 encoded data...  
...base64 encoded data)



decoded image

MODULE 6

# DNS

---

Domain Name System

1 · Network applications

2 · The Web and HTTP

3 · Email

4 · **DNS**


# Hostnames and IP addresses

- As we have seen, network routing entirely relies on **IP addresses** — numerical identifiers, e.g. `132.181.106.9`
- However, when you browse the web, you typically type a **human-friendly hostname** such as `canterbury.ac.nz`
- How does the browser know how to translate between a hostname and an IP address?



DNS

# Domain Name System — overview



# Domain Name System — overview

**DNS** is that translation mechanism — it maps human-readable **hostnames** to numerical **IP addresses**

DNS stores this mapping in a **distributed database**, spread across a hierarchy of DNS servers [RFC 1034, 1035]

DNS is an **application-layer protocol** running mainly over **UDP on port 53** (TCP port 53 for large responses and zone transfers)

Why UDP? A query and its reply each fit in one small datagram, so there is no connection setup delay, providing a quick response.



## Additional services provided by DNS

**Host aliasing** — a host may have a cumbersome **canonical hostname**, e.g. `relay1.west-coast.enterprise.com`, which is hard to remember. With host aliasing, it can also be reached via simpler **alias names**, e.g. `www.enterprise.com` and `shop.enterprise.com`; DNS maps each alias to the canonical hostname.



## Additional services provided by DNS

**Host aliasing** — a host may have a cumbersome **canonical hostname**, e.g. `relay1.west-coast.enterprise.com`, which is hard to remember. With host aliasing, it can also be reached via simpler **alias names**, e.g. `www.enterprise.com` and `shop.enterprise.com`; DNS maps each alias to the canonical hostname.

**Mail server aliasing** — the same idea for email. e.g. the mail server that handles Gmail has the cumbersome canonical hostname `gmail-smtp-in.l.google.com`. Yet Bob's email address can simply be `bob@gmail.com`: DNS uses the part after the @, `gmail.com`, to find that mail server.



# Additional services by DNS: load distribution

Many popular web services — e.g. YouTube, Google, Netflix — are hit by **millions of users** simultaneously. A **single server** cannot handle that volume.

The service is therefore distributed across **many server instances**, each with a **distinct IP address**, all sharing the **same hostname** (e.g. `youtube.com`).

## How are users assigned to these servers?

**DNS load balancing** — DNS stores all the IPs and **rotates their order** on each response, so consecutive clients get a different IP listed first.

e.g. three clients query `youtube.com` in turn:



# DNS is distributed, not centralised

Recall that DNS stores the hostname-to-IP mapping in a **distributed database**, spread across many DNS servers.

**Why not a centralised design — a single server hosting the whole database, so that every DNS query goes to it?**

## Single point of failure

If the one server crashes, DNS lookup fails for the entire Internet

## Traffic volume

Billions of DNS queries per day — no single machine can handle that load

## Distant centralised database

All queries must travel to one location — users far away face high latency on every lookup

## Maintenance

Billions of hostname records change constantly — impossible to keep one database current


# A hierarchical distributed database



## Root DNS Servers

- The top of the hierarchy: looking up a hostname from scratch **starts here**
- There are 13 root server identities (a–m.root-servers.net)
- Each is replicated: its IP address is shared by many machines worldwide, and a query goes to the nearest one


# A hierarchical distributed database



## Top-Level Domain (TLD) DNS Servers

- Handle generic domains: **.com**, **.org**, **.edu**
- ...and country-code domains: **.nz**, **.uk**
- Each is run by a dedicated registry, e.g. **Verisign** (.com), **InternetNZ** (.nz)

# A hierarchical distributed database




## Authoritative DNS Servers

- Each organisation's own DNS server
- It stores the organisation's own records — the original, **definitive** copy that every answer for its hostnames comes from
- e.g. Canterbury's authoritative server holds the IP for `www.canterbury.ac.nz`

# Local DNS server and DNS caching

Every ISP or institution operates a **local DNS server** (also called a *default name server*).

It is **not** part of the root / TLD / authoritative hierarchy — it sits **separately** from that.




**Local DNS server** dns.canterbury.ac.nz



# Local DNS server and DNS caching

When you join the network, your device is automatically configured to use one — e.g. the University of Canterbury runs its own local DNS server for all campus users.




**Local DNS server** dns.canterbury.ac.nz



Your computer


# Local DNS server and DNS caching

When your machine makes a DNS lookup, the query goes **first to the local DNS server**.



# Local DNS server and DNS caching

**If the answer is cached:** the local DNS server replies **immediately** — no hierarchy lookup needed.




# Local DNS server and DNS caching

**If not cached:** the local server **forwards the query into the hierarchy.**

The hierarchy replies with the IP address.

The local server **relays the reply** back to you,




# Local DNS server and DNS caching

**If not cached:** the local server **forwards the query into the hierarchy**.


The hierarchy replies with the IP address.

The local server **relays the reply** back to you,  
and **caches the hostname-to-IP mapping** for future queries.





# Local DNS server and DNS caching

Cache entries have a **TTL (time-to-live)** and expire after some time (typically hours). TLD server addresses are routinely cached in local DNS servers — so **root servers are bypassed for most everyday queries**.




# Resolving a hostname: the full journey

**Step 1.** A client wants to resolve **csse.canterbury.ac.nz**. It sends a DNS query to its **local DNS server**.




# Resolving a hostname: the full journey

**Step 2.** The local DNS doesn't have the answer cached. It contacts a **root DNS server**.




# Resolving a hostname: the full journey

**Step 3.** The root server returns a **referral**: "I don't know — try a **.nz TLD** server" and gives its address.




# Resolving a hostname: the full journey

**Step 4.** The local DNS follows the referral and contacts the **.nz TLD server**.




# Resolving a hostname: the full journey

**Step 5.** The .nz TLD returns another **referral**:  
"Try **dns.canterbury.ac.nz** — the  
authoritative server for canterbury.ac.nz."




# Resolving a hostname: the full journey

**Step 6.** The local DNS contacts **dns.canterbury.ac.nz** — the authoritative name server for the canterbury.ac.nz zone.




# Resolving a hostname: the full journey

**Step 7.** The authoritative server has the answer! It replies: "csse.canterbury.ac.nz = 132.181.2.21"



# Resolving a hostname: the full journey


**Step 8.** The local DNS returns **132.181.2.21** to the client — and **caches** the mapping for future queries.



# Recursive vs iterative queries

## Iterative query


- The contacted server does **not** query other servers for you: it replies only with what it knows itself.
- If it doesn't know the answer, it returns a **referral**: "I don't know — try that server."
- The authoritative server *does* know, so its reply to the same kind of query is the **answer**.



# Recursive vs iterative queries

## Recursive query


- The contacted server takes **full responsibility**:
- it resolves the name itself, querying the next servers on your behalf,
- and returns the **final answer**. e.g. client → local DNS server.



# Recursive vs iterative queries

## Can we go fully recursive?

- In principle, yes: root asks the TLD server, which asks the authoritative server,
- and the answer flows back along the chain.
- In practice, root and TLD servers **refuse recursion**: too much load.



# DNS Resource Records

DNS’s distributed database stores its mappings as **Resource Records (RRs)**: each DNS server keeps a table of records, all in the same four-field format.

**name**  
e.g. a hostname

**value**  
e.g. an IP address

**type**  
A, NS, CNAME, MX

**TTL**  
time-to-live (s)

Type A

**name** = hostname · **value** = IPv4 address

| NAME                   | VALUE                  | TYPE  | TTL (s) |
|------------------------|------------------------|-------|---------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | 3600    |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | 3600    |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | 3600    |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | 86400   |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | 3600    |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | 3600    |

# DNS Resource Records

DNS's distributed database stores its mappings as **Resource Records (RRs)**: each DNS server keeps a table of records, all in the same four-field format.

**name**  
e.g. a hostname

**value**  
e.g. an IP address

**type**  
A, NS, CNAME, MX

**TTL**  
time-to-live (s)

Type NS

**name** = domain · **value** = hostname of the authoritative DNS server for that domain

*"I don't know — but here's the name of the server that does."*

| NAME                   | VALUE                  | TYPE  | TTL (s) |
|------------------------|------------------------|-------|---------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | 3600    |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | 3600    |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | 3600    |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | 86400   |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | 3600    |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | 3600    |

# DNS Resource Records

DNS’s distributed database stores its mappings as **Resource Records (RRs)**: each DNS server keeps a table of records, all in the same four-field format.

**name**  
e.g. a hostname

**value**  
e.g. an IP address

**type**  
A, NS, CNAME, MX

**TTL**  
time-to-live (s)

**Type CNAME**

**name** = alias hostname · **value** = canonical (real) hostname it maps to

| NAME                   | VALUE                  | TYPE  | TTL (s) |
|------------------------|------------------------|-------|---------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | 3600    |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | 3600    |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | 3600    |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | 86400   |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | 3600    |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | 3600    |

# DNS Resource Records

DNS's distributed database stores its mappings as **Resource Records (RRs)**: each DNS server keeps a table of records, all in the same four-field format.

**name**  
e.g. a hostname

**value**  
e.g. an IP address

**type**  
A, NS, CNAME, MX

**TTL**  
time-to-live (s)

## Type MX


**name** = domain · **value** = hostname of the mail server for that domain

Unlike CNAME, it only redirects mail: a sending mail server (**SMTP**) looks it up to find where to deliver a domain's email.

| NAME                   | VALUE                  | TYPE  | TTL (s) |
|------------------------|------------------------|-------|---------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | 3600    |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | 3600    |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | 3600    |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | 86400   |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | 3600    |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | 3600    |

# DNS protocol message format


DNS uses the **same message format** for both **queries** and **replies** — a fixed 12-byte header followed by four variable-length sections.



# DNS protocol message format

## Identification


16-bit number assigned by the client. The reply copies this same value, so the client can match each reply to its original query.



# DNS protocol message format

## Flags


- Query (0) or reply (1)
- Recursion desired (RD) — client requests full recursive resolution
- Recursion available (RA) — server indicates it supports recursion



# DNS protocol message format

## Count fields

Four 16-bit counters — one per section below — telling how many records appear in each variable-length section.




# DNS protocol message format

## Questions

The name and record type being queried. e.g. this Type A question asks for the IP address of the domain facebook.com.

(facebook.com, A)




# DNS protocol message format

## Answers

Resource Records (RRs) that directly answer the query.

```
(facebook.com, 157.240.22.35, A,
TTL=300)
```




# DNS protocol message format

## Authority

NS records pointing to authoritative servers —  
returned in referral responses.

(facebook.com, a.ns.facebook.com, NS)




# DNS protocol message format

## Additional

'Glue' Resource Records (RRs) accompanying the authority section — so you can reach the name server without an extra lookup.

(a.ns.facebook.com, 129.134.30.12, A)



# Query & response: a live example

## QUERY client → local DNS · recursive

Transaction ID 0x6492  
Flags 0x0100 (query, RD=1)  
Questions 1 Answer RRs 0  
Authority RRs 0 Additional RRs 0

### QUESTIONS

www.facebook.com A

### The question

The query asks one question: what is the IP address (type **A**) of `www.facebook.com`?

## RESPONSE local DNS → client

Transaction ID 0x6492  
Flags 0x8180 (reply, RA=1, no error)  
Questions 1 Answer RRs 2  
Authority RRs 4 Additional RRs 8

### QUESTIONS

www.facebook.com A

### ANSWERS

www.facebook.com CNAME  
star-mini.c10r.facebook.com  
star-mini.c10r.facebook.com A  
31.13.78.35

### AUTHORITY

c10r.facebook.com NS  
d.ns.c10r.facebook.com

### ADDITIONAL

d.ns.c10r.facebook.com A  
185.89.219.11

# Query & response: a live example

## QUERY client → local DNS · recursive

Transaction ID 0x6492  
Flags 0x0100 (query, RD=1)  
Questions 1 Answer RRs 0  
Authority RRs 0 Additional RRs 0

### QUESTIONS

www.facebook.com A

### The answer

The response answers with **two records**.

## RESPONSE local DNS → client

Transaction ID 0x6492  
Flags 0x8180 (reply, RA=1, no error)  
Questions 1 Answer RRs 2  
Authority RRs 4 Additional RRs 8

### QUESTIONS

www.facebook.com A

### ANSWERS

www.facebook.com CNAME  
star-mini.c10r.facebook.com  
star-mini.c10r.facebook.com A  
31.13.78.35

### AUTHORITY

c10r.facebook.com NS  
d.ns.c10r.facebook.com

### ADDITIONAL

d.ns.c10r.facebook.com A  
185.89.219.11

# Query & response: a live example

## QUERY client → local DNS · recursive

Transaction ID 0x6492  
Flags 0x0100 (query, RD=1)  
Questions 1 Answer RRs 0  
Authority RRs 0 Additional RRs 0

### QUESTIONS

www.facebook.com A

### First: a CNAME

www.facebook.com is an **alias**; its canonical name is star-mini.c10r.facebook.com.

## RESPONSE local DNS → client

Transaction ID 0x6492  
Flags 0x8180 (reply, RA=1, no error)  
Questions 1 Answer RRs 2  
Authority RRs 4 Additional RRs 8

### QUESTIONS

www.facebook.com A

### ANSWERS

www.facebook.com CNAME  
star-mini.c10r.facebook.com  
star-mini.c10r.facebook.com A  
31.13.78.35

### AUTHORITY

c10r.facebook.com NS  
d.ns.c10r.facebook.com

### ADDITIONAL

d.ns.c10r.facebook.com A  
185.89.219.11

# Query & response: a live example

## QUERY client → local DNS · recursive

Transaction ID 0x6492  
Flags 0x0100 (query, RD=1)  
Questions 1 Answer RRs 0  
Authority RRs 0 Additional RRs 0

### QUESTIONS

www.facebook.com A

### Then: an A record

The canonical name's IP address, 31.13.78.35 — the answer we wanted.

## RESPONSE local DNS → client

Transaction ID 0x6492  
Flags 0x8180 (reply, RA=1, no error)  
Questions 1 Answer RRs 2  
Authority RRs 4 Additional RRs 8

### QUESTIONS

www.facebook.com A

### ANSWERS

www.facebook.com CNAME  
star-mini.c10r.facebook.com  
star-mini.c10r.facebook.com A  
31.13.78.35

### AUTHORITY

c10r.facebook.com NS  
d.ns.c10r.facebook.com

### ADDITIONAL

d.ns.c10r.facebook.com A  
185.89.219.11

# Query & response: a live example

## QUERY client → local DNS · recursive

Transaction ID 0x6492  
Flags 0x0100 (query, RD=1)  
Questions 1 Answer RRs 0  
Authority RRs 0 Additional RRs 0

### QUESTIONS

www.facebook.com A

### Authority

An **NS** record naming an authoritative name server for c10r.facebook.com.

## RESPONSE local DNS → client

Transaction ID 0x6492  
Flags 0x8180 (reply, RA=1, no error)  
Questions 1 Answer RRs 2  
Authority RRs 4 Additional RRs 8

### QUESTIONS

www.facebook.com A

### ANSWERS

www.facebook.com CNAME  
star-mini.c10r.facebook.com  
star-mini.c10r.facebook.com A  
31.13.78.35

### AUTHORITY

c10r.facebook.com NS  
d.ns.c10r.facebook.com

### ADDITIONAL

d.ns.c10r.facebook.com A  
185.89.219.11

# Query & response: a live example

## QUERY client → local DNS · recursive

Transaction ID 0x6492  
Flags 0x0100 (query, RD=1)  
Questions 1 Answer RRs 0  
Authority RRs 0 Additional RRs 0

### QUESTIONS

www.facebook.com A

### Additional

The **IP address** of that name server, so it can be reached directly.

## RESPONSE local DNS → client

Transaction ID 0x6492  
Flags 0x8180 (reply, RA=1, no error)  
Questions 1 Answer RRs 2  
Authority RRs 4 Additional RRs 8

### QUESTIONS

www.facebook.com A

### ANSWERS

www.facebook.com CNAME  
star-mini.c10r.facebook.com  
star-mini.c10r.facebook.com A  
31.13.78.35

### AUTHORITY

c10r.facebook.com NS  
d.ns.c10r.facebook.com

### ADDITIONAL

d.ns.c10r.facebook.com A  
185.89.219.11

# DNS — key takeaways

- **Main purpose** — DNS maps human-readable hostnames to IP addresses.
- **Hierarchy** — DNS is distributed: no single server holds all records.  
Responsibility is delegated across Root, TLD, and Authoritative servers.
- **Traversal pattern** — the client queries the local DNS resolver, which then follows referrals Root → TLD → Authoritative.  
Results are cached locally by TTL to avoid repeated lookups.

## References

James F. Kurose & Keith W. Ross, *Computer Networking: A Top-Down Approach*, Addison Wesley, 3<sup>rd</sup> edition, 2005.

Chapter 2 Application Layer (2.1, 2.2, 2.4, 2.5)

William Stallings, *Data and Computer Communications*, Pearson, 10<sup>th</sup> edition, 2013.

Chapter 24 Electronic Mail, DNS, and HTTP (24.1–24.3)

Andrew S. Tanenbaum & David J. Wetherall, *Computer Networks*, Prentice Hall, 5<sup>th</sup> edition, 2010.

Chapter 7 The Application Layer (7.1–7.3)

# Acknowledgements

These lecture notes are developed based on slides and notes from:

Barry Wu's lecture notes for COSC264, University of Canterbury

Dr Mengmeng Ge's lecture notes for COSC264, University of Canterbury, 2024

Assoc Prof Dan Kim's lecture notes for COSC264, University of Canterbury, 2017

Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University, 2005