

COSC264 · INTRODUCTION TO COMPUTER NETWORKS AND THE INTERNET · TERM 4 · MODULE 6

# Application Protocols: the Web, Email and DNS

1 · Network applications

2 · The Web and HTTP

3 · Email

4 · DNS

Module 6 – Application Protocols

[cosc264.up.railway.app](https://cosc264.up.railway.app)

2 Oct 2026 · 1 / 50

OUTLINE

# Module 6 — roadmap

**General concepts about network applications**

**One network application: the Web**

**HTTP: the protocol that supports the Web**

**Email**

**DNS: the Domain Name System**

MODULE 6

# Network Applications

---

What an application needs from the layers beneath it

1 • Network applications

2 • The Web and HTTP

3 • Email

4 • DNS

# What is a network application?

**Programs that run on different end systems** and communicate with each other over the network.

## Web browsing

Google Chrome, Mozilla Firefox,  
Microsoft Edge

## Email

Gmail, Outlook

## Instant messaging

Slack, WhatsApp, Facebook  
Messenger

## Streaming services

YouTube, Netflix

## Social media

Instagram, TikTok

## AI assistants

Claude, ChatGPT, Gemini

# Application-layer protocols

An **application-layer protocol** is one piece of a network application — but a big piece.

It defines how the application's processes, running on different end systems, **pass messages to each other**: which messages exist, their format, and when and how to respond.

This is what lets programs on different machines work together — e.g. any web browser can fetch a page from any web server, whoever wrote each program.

| Network application | Application-layer protocol |
|---------------------|----------------------------|
| Web application     | HTTP                       |
| Email               | SMTP                       |

## Choosing a transport protocol

- Application-layer protocols are controlled by the **application developer**; transport-layer protocols, however, are built into the **OS**.
- As a developer you can only choose between the transport protocols the OS provides — so pick the one that best fits what your application-layer protocol needs.

| Property               | Example applications                                | Transport protocol                                      |
|------------------------|-----------------------------------------------------|---------------------------------------------------------|
| Reliable data transfer | Email, IM, file transfer, financial apps            | <b>TCP</b> preferred<br>UDP unsuitable                  |
| Low delay              | Internet telephony, conferencing, multiplayer games | <b>UDP</b> preferred<br>TCP retransmissions add latency |

MODULE 6

# The Web and HTTP

---

The application-layer protocol behind every page you load

1 · Network applications

2 · **The Web and HTTP**

3 · Email

4 · DNS

# The World Wide Web

## THE WEB

The **World Wide Web** (aka Web) is a **network application**: a global collection of documents and other resources shared over the Internet — like a large file storage system.

Each object on the Web (e.g., an image, an HTML file) is addressable by a **URL**:

```
http://www.someschool.edu/someDept/pic.gif  
└─ host name ─┐ └─ path name ─┘
```


## HTTP

# HTTP: the protocol behind the Web

### HTTP

The **Hypertext Transfer Protocol** is the **application-layer protocol** that supports the Web — it defines how clients request objects by URL and how servers respond.

HTTP follows a **request–response** pattern: the client sends an **HTTP request** for an object, and the server answers with an **HTTP response** carrying it.



## HTTP

# How HTTP works

HTTP uses **TCP** as its transport protocol — so before any HTTP message is sent, a TCP connection must be established first.



### STATE LESS

HTTP is **stateless** — the server maintains no information about past client requests. Each request is handled independently.

# HTTP connections

## Non-persistent HTTP

- At most **one object** is sent over a TCP connection — connection is then closed
- HTTP/1.0 uses non-persistent HTTP

## Persistent HTTP


- **Multiple objects** can be sent over a single TCP connection
- HTTP/1.1, HTTP/2, and HTTP/3 use persistent HTTP

| Protocol | Published | Connection type   | Status   |
|----------|-----------|-------------------|----------|
| HTTP/1.0 | 1996      | Non-persistent    | Legacy   |
| HTTP/1.1 | 1997      | Persistent        | Standard |
| HTTP/2   | 2015      | Persistent        | Standard |
| HTTP/3   | 2022      | Persistent (QUIC) | Standard |

# Non-persistent HTTP in action

Suppose a client wants to retrieve a file from: `http://www.someSchool.edu/someDepartment/home.index`

1. Client establishes a **TCP connection** with the server
2. Client sends an **HTTP GET** request — path: `/someDepartment/home.index`
3. Server retrieves the HTML file, wraps it in an **HTTP response**, and sends it back
4. Server signals TCP to **close the connection**



## Response time: non-persistent HTTP

### RESPONSE TIME

The time to retrieve a single object using HTTP — from the client's request until the object is received

**Q:** What is the response time for non-persistent HTTP?

**A:** 2 RTTs + file transmission time

### RTT

Recall: **round-trip time** — the time for a small packet to travel from client to server and back

**First RTT:** the TCP handshake — client sends SYN, server replies SYN-ACK

**Second RTT:** client sends HTTP GET, server sends back the response

With many referenced objects, this cost multiplies — a significant burden on the server



## Persistent HTTP in action

Client and server establish a TCP connection and exchange HTTP messages

Server **keeps the TCP connection open** after sending the first HTTP response

Subsequent HTTP messages between the same client and server use the **same TCP connection** — no new handshake needed

### Response time on one persistent connection:

1 RTT for the initial TCP handshake, then 1 RTT for each HTTP request.

So for  $n$  requests: **total**  $\approx (1 + n)$  RTT (plus transmission time).

Amortised over the  $n$  objects:  $(1 + n)/n = (1/n + 1)$  RTT. As  $n$  becomes large,  $1/n \rightarrow 0$ , so this approaches **~1 RTT per object**.

Can we do even better than 1 RTT per object?




HTTP

## Improving HTTP efficiency further with pipelining

**WITHOUT  
PIPELINING**

Client issues a new request **only after** receiving the previous response —  
**1 RTT per file**



# Improving HTTP efficiency further with pipelining

## WITHOUT PIPELINING


Client issues a new request **only after** receiving the previous response — **1 RTT per file**

## WITH PIPELINING

Client sends requests **as soon as it needs to retrieve an object**, without waiting for responses — e.g. in this example we pipeline 3 HTTP requests all together

**Benefit:** pipelining 3 requests → response time per file  $\approx 1 \text{ RTT} / 3$

- Don't confuse this with TCP-level pipelining.
- **HTTP pipelining** operates at the **application layer**: the client sends several **HTTP requests** without waiting for each response.
- **TCP pipelining** operates at the **transport layer**: the sender keeps several **segments** in flight within its window. TCP cuts each HTTP message into segments, so these may be pieces of one HTTP request, or of several.



# HTTP request message format

Method — what action to perform

| Method        | Action                          |
|---------------|---------------------------------|
| <b>GET</b>    | Retrieve a resource             |
| <b>HEAD</b>   | Retrieve headers only (no body) |
| <b>POST</b>   | Submit data to a resource       |
| <b>PUT</b>    | Store / replace a resource      |
| <b>DELETE</b> | Remove a resource               |

Request line

method

URL

version

cr lf

Header lines

field name:

value

cr lf

field name:

value

cr lf

Blank line

cr lf

Entity body

*entity body*

# HTTP request message format

**URL:** path to the requested resource — e.g.

`/someDept/pic.gif`

**HTTP Version:** identifies the protocol in use — HTTP/1.0 or HTTP/1.1

**Request line**

method

URL

version

cr lf

**Header lines**

field name:

value

cr lf

field name:

value

cr lf

**Blank line**

cr lf

**Entity body**

*entity body*

# HTTP request message format

## Header lines

- Zero or more `field-name: value` pairs, one per line
- Common examples:

```
Host: www.example.com
```

Mandatory in HTTP/1.1 — identifies the target server

```
Content-Type: text/html; charset=UTF-8
```

Format of the message body

```
Connection: close
```

Close the TCP connection after this response

```
Accept-Language: da, en-gb;q=0.8,  
en;q=0.7
```

Preferred languages, weighted by q (default 1): Danish → British English → any English

## Request line

method URL version cr lf

## Header lines

field name: value cr lf

field name: value cr lf

## Blank line


cr lf

## Entity body

*entity body*

# HTTP request message format

**Entity body:** empty for GET and HEAD; carries the payload for POST and PUT (e.g. form data or JSON)



HTTP

## HTTP request: an example

```
GET /somedir/page.html HTTP/1.1
```

```
Host: www.someschool.edu
```

```
User-agent: Mozilla/4.0
```

```
Connection: close
```

```
Accept-language: fr
```

-----  
*(empty entity body)*

**Method:** GET

**URL:** /somedir/page.html

**HTTP version:** HTTP/1.1

**Header field:** Host: www.someschool.edu

**Header field:** User-agent: Mozilla/4.0

**Header field:** Connection: close

**Header field:** Accept-language: fr

**Body:** (empty)

# HTTP response message format

## Version

The HTTP version the server is using — e.g. HTTP/1.1, just as in the request line.

Status line

version status code phrase cr lf

Header lines

field name: value cr lf

field name: value cr lf

Blank line

cr lf

Entity body

*entity body*

# HTTP response message format

## Status code & phrase

- **Status code:** 3-digit number indicating the result of the request
- **Phrase:** human-readable text accompanying the code — e.g. OK, Not Found

## Status line

version — status code — phrase — cr lf

## Header lines

field name: — value — cr lf

field name: — value — cr lf

## Blank line

cr lf

## Entity body

*entity body*

# HTTP response message format

## Common examples

| Code | Phrase                     |
|------|----------------------------|
| 200  | OK                         |
| 301  | Moved Permanently          |
| 400  | Bad Request                |
| 404  | Not Found                  |
| 505  | HTTP Version Not Supported |


### Status line




### Header lines



### Blank line




### Entity body



# HTTP response message format

## Header lines, blank line, entity body

Same structure as in the request: `field-name: value`  
header lines, a blank line (a lone CRLF), then the body —  
here, the **requested object** itself.



# HTTP response: an example

```
HTTP/1.1 200 OK
Connection: close
Date: Thu, 06 Aug 1998 12:00:15 GMT
Server: Apache/1.3.0 (Unix)
Last-Modified: Mon, 22 Jun 1998 ...
Content-Length: 6821
Content-Type: text/html
```

```
<!DOCTYPE html>
<html><head><title>Example</title></head>
<body><h1>Hello!</h1>
<script>console.log("hi");</script>
</body></html>
```

**Version:** HTTP/1.1

**Status code:** 200 — success

**Phrase:** OK

**Header:** Connection: close — close after this response

**Header:** Date — when the response was sent

**Header:** Server — the server software

**Header:** Last-Modified — when the object last changed

**Header:** Content-Length — body size in bytes

**Header:** Content-Type — the body is HTML

**Body:** the requested HTML file (with embedded JS)

# Cookies


HTTP is **stateless** — every request is independent; the server retains no memory of past requests.

So the server does not remember your earlier conversations with it. Yet many applications *need* that memory: when you refresh an online shop's page, the items in your cart should not just disappear. **How does the server know what is in your cart?**

The answer: **cookies** — here's how it works.

1. A **Set-Cookie** header line in the **HTTP response** — the server sends the client a unique ID on first contact
2. A **Cookie** header line in subsequent **HTTP requests** — the client echoes the ID back on every future request
3. A **cookie file** stored on the client, managed by the browser — persists across sessions
4. A **back-end database** on the server, keyed by the cookie ID — stores per-user state

# Cookies: an example [RFC 6265]




The client sends an HTTP request to the host

**Set-Cookie: id=1678** — Server assigns a unique ID and sends it to the client

The cookie ID is stored as an entry in the back-end database, and in the client's browser

# Cookies: an example [RFC 6265]



**Cookie: id=1678** — Client attaches the ID on every future request to the same server

Server uses the ID to look up the user's state in its database

Server can then generate a **tailored response** — e.g. a personalised page, a pre-filled shopping cart — for the client's request

HTTP

## The problem: every request goes to the origin


Picture a university: thousands of users fetch the same popular pages and images, and every request travels across the Internet to the **origin server** that hosts the object.

Every one of those requests lands on the same origin server, through the same congested links — enough of them can **overwhelm** it.

**How can we avoid this?**

# Web caching

A **web cache** (also called a **proxy server**) is a network entity that satisfies HTTP requests on behalf of an origin server — storing copies of recently requested objects locally.



## Acts as a server

Receives HTTP requests from clients and returns responses — just like an origin server would

## Acts as a client

When the object is not cached, it issues its own HTTP request to the origin server to fetch it

HTTP

# What web caching buys us


## **Quicker responses**

A cached object is served from a nearby cache, with no round trip to a distant origin server — everyone gets their object sooner.

## **Less congestion**

Requests answered by the cache never reach the origin server or cross the institution's access link — so neither is overwhelmed.

# Conditional GET



**Goal:** the proxy server/cache wants to avoid re-fetching an object if its cached copy is still up-to-date

The proxy sends the following HTTP GET request:

```
GET /index.html HTTP/1.1
```

```
Host: www.example.com
```

```
If-modified-since: Mon, 10 Aug 2026 09:00:00  
GMT
```

If the object is **unchanged**, server replies:

```
304 Not Modified
```

No object in the body — cache serves its stored copy directly

Object **changed**: server replies 200 OK with the updated object in the body

MODULE 6

# Email

---

Electronic Mail

1 · Network applications

2 · The Web and HTTP

**3 · Email**

4 · DNS


## OUTLINE

# Email — roadmap

### Electronic Mail (Email)

- *Mail transfer protocol:* **SMTP**
- *Mail access protocols:* **POP3, IMAP**
- *Message format:* **MIME**

# An overview of the email system



**User agent (UA)** — the app used to read, compose and send email (e.g. Outlook, Gmail, Apple Mail)


**Mail server** — a dedicated host that handles email on behalf of its users. It stores a **mailbox** per recipient for incoming mails and a **send queue** of outgoing mails

**SMTP** — Simple Mail Transfer Protocol; used to send mail, from the sender's UA to its mail server and between mail servers across the internet

**Mail access protocol (MAP)** — used by UAs to retrieve messages from their mail server; POP3 or IMAP


## SMTP in action

SMTP moves messages between mail servers. It consists of a **client side** and a **server side**.  
Every mail server runs both an **SMTP client** (when sending) and an **SMTP server** (when receiving).



- ① Alice uses her UA to **compose** the message
- ② Alice's UA **sends** the message to her mail server (also via SMTP)
- ③ The **SMTP client** on Alice's server opens a TCP connection to Bob's mail server
- ④ Alice's message is **transferred over the TCP connection**
- ⑤ The **SMTP server** on Bob's server receives the message and places it in Bob's **mailbox**
- ⑥ Bob invokes his UA to **read** the message

## Retrieving email from the server



**Mail access protocols** retrieve email from the mail server to the recipient's user agent

**POP3** — Post Office Protocol version 3 [RFC 1939]  
Downloads messages to the local device. Simple and stateless between sessions

**IMAP** — Internet Message Access Protocol version 4 [RFC 9051]  
Keeps messages on the server; supports folders and sync across multiple devices

# Accessing email via a browser

**Web-based email is accessed through a web browser**

Uses **HTTP** for communication between the user's browser and the mail server

**Web-based email services**

Examples: **Gmail**, **Outlook** (web)

IMAP


# Internet Message Access Protocol

IMAP **keeps messages on the server** and replicates copies to the client

Users can **organise messages in folders** — folder structure is stored on the server, accessible from any device

IMAP maintains **user state information** across sessions (e.g. folder names, message-to-folder mappings)

## How POP3 works



### Notable limitations:

- Does not synchronise emails across multiple devices
- Does not support server-side folder management — users cannot organise emails into folders on the server

# Structure of an email message

An email message is **ASCII text** — **header fields** followed (optionally) by a **body**

```
Date: Mon, 10 Aug 2026 09:00:00
GMT
From: alice@crepes.fr
To: bob@hamburger.edu
Subject: Picture of yummy crepe.
Cc: charlie@example.com
```

Header lines

---

```
Hello Bob, here is the picture
you asked for. Enjoy!
```

Body

## MIME

# Multipurpose Internet Mail Extensions

**Problem:** standard email is ASCII-only — it cannot natively carry images, attachments, or non-ASCII characters (e.g. Chinese, Arabic)

MIME extends the message format with additional header fields that describe and encode non-ASCII content — telling the receiver how to decode the body

e.g. **base64** encoding converts binary data into printable ASCII text, which can be used for image attachments

From: alice@crepes.fr

To: bob@hamburger.edu

Subject: Picture of yummy crepe.

**MIME-Version: 1.0**

**Content-Transfer-Encoding: base64**

**Content-Type: image/jpeg**

---

(base64 encoded data...

...base64 encoded data)

MODULE 6

# DNS

---

Domain Name System

1 · Network applications

2 · The Web and HTTP

3 · Email

4 · **DNS**

# Domain Name System — overview

- As we have seen, network routing entirely relies on **IP addresses** — numerical identifiers
- However, when you browse the web, you typically type a **human-friendly hostname** such as `canterbury.ac.nz`
- How does the browser know how to translate between a hostname and an IP address?




# Domain Name System — overview

- As we have seen, network routing entirely relies on **IP addresses** — numerical identifiers
- However, when you browse the web, you typically type a **human-friendly hostname** such as `canterbury.ac.nz`
- How does the browser know how to translate between a hostname and an IP address?

**DNS** is that translation mechanism — it maps human-readable **hostnames** to numerical **IP addresses**

DNS implements a **distributed database** — a global hostname-to-IP mapping distributed across a hierarchy of DNS servers [RFC 1034, 1035]

DNS is an **application-layer protocol** running mainly over **UDP on port 53** (TCP port 53 for large responses and zone transfers)




## Additional services provided by DNS

**Hostname-to-IP translation** — the primary service, already seen; commonly used by HTTP, SMTP, and other application-layer protocols

**Host aliasing** — a host with a complex **canonical hostname** can be reached via simpler **alias names**

e.g. `www.enterprise.com` and `shop.enterprise.com` both resolve to the canonical `relay1.west-coast.enterprise.com`

**Mail server aliasing** — the same idea for email; a friendly address like `bob@enterprise.com` can map (via an MX record) to the canonical mail server hostname `relay1.west-coast.enterprise.com`




## Additional services by DNS: load distribution

Many popular web services — e.g. YouTube, Google, Netflix — are hit by **millions of users** simultaneously. A **single server** cannot handle that volume.

The service is therefore distributed across **many server instances**, each with a **distinct IP address**, all sharing the **same hostname** (e.g. `youtube.com`).

How are users assigned to these servers? **DNS load balancing** — DNS stores all the IPs and **rotates their order** on each response, so consecutive clients get a different IP listed first.



# Why not a centralised design?

A centralised design would use a **single server** that hosts the translation database — every DNS query goes to that server.

## Single point of failure

If the one server crashes, DNS lookup fails for the entire Internet

## Traffic volume

Billions of DNS queries per day — no single machine can handle that load


## Distant centralised database

All queries must travel to one location — users far away face high latency on every lookup

## Maintenance

Billions of hostname records change constantly — impossible to keep one database current


# A hierarchical distributed database



## Root DNS Servers

13 root server identities (a–m.root-servers.net) — each replicated via anycast, giving **over 1000 physical instances** worldwide

# A hierarchical distributed database




## TLD DNS Servers

Handle **.com**, **.org**, **.edu** and country-code domains (.nz, .uk). Maintained by dedicated registries — e.g.

**Verisign** (.com), **Public Interest Registry** (.org), **EDUCAUSE** (.edu), **InternetNZ** (.nz).

# A hierarchical distributed database



## Authoritative DNS Servers

Each organisation's own DNS server — the **definitive source** for its hostname-to-IP mappings.  
e.g. Canterbury's authoritative server holds the IP for `www.canterbury.ac.nz`

# Local DNS server and DNS caching

Every ISP or institution operates a **local DNS server** (also called a *default name server*). It is **not** part of the root / TLD / authoritative hierarchy — it sits **separately** from that.

When you join the network, your device is automatically configured to use one — e.g. the University of Canterbury runs its own local DNS server for all campus users.

**DNS Hierarchy**  
Root · TLD · Authoritative  
(separate from local DNS)

**Local DNS server**  
e.g. dns.canterbury.ac.nz  
(cache)

Your computer


# Local DNS server and DNS caching

When your machine makes a DNS lookup, the query goes **first to the local DNS server**.

**If the answer is cached:** the local DNS server replies **immediately** — no hierarchy lookup needed.


**If not cached:** the local server **forwards the query into the hierarchy**, relays the reply back to you, and **caches the hostname-to-IP mapping** in local memory for future queries.

Cache entries have a **TTL (time-to-live)** and expire after some time (typically hours). TLD server addresses are routinely cached in local DNS servers — so **root servers are bypassed for most everyday queries**.




# Resolving a hostname: the full journey

**Step 1.** A client wants to resolve **csse.canterbury.ac.nz**. It sends a DNS query to its **local DNS server**.




# Resolving a hostname: the full journey

**Step 2.** The local DNS doesn't have the answer cached. It contacts a **root DNS server**.




# Resolving a hostname: the full journey

**Step 3.** The root server returns a **referral**: "I don't know — try a **.nz TLD** server" and gives its address.




# Resolving a hostname: the full journey

**Step 4.** The local DNS follows the referral and contacts the **.nz TLD server**.




# Resolving a hostname: the full journey

**Step 5.** The .nz TLD returns another **referral**:  
"Try **dns.canterbury.ac.nz** — the  
authoritative server for canterbury.ac.nz."




# Resolving a hostname: the full journey

**Step 6.** The local DNS contacts **dns.canterbury.ac.nz** — the authoritative name server for the canterbury.ac.nz zone.




# Resolving a hostname: the full journey

**Step 7.** The authoritative server has the answer! It replies: "csse.canterbury.ac.nz = 132.181.2.21"



# Resolving a hostname: the full journey

**Step 8.** The local DNS returns **132.181.2.21** to the client — and **caches** the mapping for future queries.




# Recursive vs iterative queries

## Iterative query

The contacted server does **not** resolve the name for you. It returns a **referral** — "I don't know, try that server."

## Recursive query

The contacted server takes **full responsibility** — it resolves the name itself (querying the next server in the chain) and returns the **final answer**.



# Recursive vs iterative queries

## Iterative query


The contacted server does **not** resolve the name for you. It returns a **referral** — "I don't know, try that server."

## Recursive query

The contacted server takes **full responsibility** — it resolves the name itself (querying the next server in the chain) and returns the **final answer**.

## Can we go fully recursive?

In principle, yes — root asks TLD, which asks auth. In practice, root and TLD servers refuse recursion (too much load), so only client → local DNS is recursive.



# DNS Resource Records

DNS distributes data as **Resource Records (RRs)**. Each record has four fields:

(name, value, type, TTL – Time-To-Live)

| NAME                   | VALUE                  | TYPE  | TTL  |
|------------------------|------------------------|-------|------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | XXXX |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | XXXX |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | XXXX |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | XXXX |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | XXXX |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | XXXX |

# DNS Resource Records

Type A

**name** = hostname · **value** = IPv4 address

| NAME                   | VALUE                  | TYPE  | TTL  |
|------------------------|------------------------|-------|------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | XXXX |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | XXXX |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | XXXX |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | XXXX |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | XXXX |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | XXXX |

# DNS Resource Records

Type NS

**name** = domain · **value** = hostname of the authoritative DNS server for that domain

*"I don't know — but here's the name of the server that does."*

| NAME                   | VALUE                  | TYPE  | TTL  |
|------------------------|------------------------|-------|------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | XXXX |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | XXXX |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | XXXX |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | XXXX |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | XXXX |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | XXXX |

# DNS Resource Records

Type CNAME

**name** = alias hostname · **value** = canonical (real) hostname it maps to

| NAME                   | VALUE                  | TYPE  | TTL  |
|------------------------|------------------------|-------|------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | XXXX |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | XXXX |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | XXXX |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | XXXX |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | XXXX |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | XXXX |

# DNS Resource Records

Type MX

**name** = domain · **value** = hostname of the mail server for that domain

| NAME                   | VALUE                  | TYPE  | TTL  |
|------------------------|------------------------|-------|------|
| web.canterbury.ac.nz   | 132.181.106.9          | A     | XXXX |
| itdns.canterbury.ac.nz | 132.181.2.225          | A     | XXXX |
| mail.canterbury.ac.nz  | 132.181.3.225          | A     | XXXX |
| canterbury.ac.nz       | itdns.canterbury.ac.nz | NS    | XXXX |
| www.canterbury.ac.nz   | web.canterbury.ac.nz   | CNAME | XXXX |
| canterbury.ac.nz       | mail.canterbury.ac.nz  | MX    | XXXX |

# DNS protocol message format

DNS uses the **same message format** for both **queries** and **replies** — a fixed 12-byte header followed by four variable-length sections.

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# DNS protocol message format

**Identification**

16-bit number assigned by the client. The reply copies this same value, so the client can match each reply to its original query.

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# DNS protocol message format

**Flags**

- Query (0) or reply (1)
- Recursion desired — client requests full recursive resolution
- Recursion available — server indicates it supports recursion

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# DNS protocol message format

Count fields

Four 16-bit counters — one per section below — telling how many records appear in each variable-length section.

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# DNS protocol message format

Questions

The name and record type being queried. e.g. this  
Type A question asks for the IP address of the domain  
facebook.com.

(facebook.com, A)

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# DNS protocol message format

Answers

Resource Records (RRs) that directly answer the query.

( facebook.com, 157.240.22.35, A, TTL=300 )

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# DNS protocol message format

**Authority**

NS records pointing to authoritative servers —  
returned in referral responses.

(facebook.com, a.ns.facebook.com, NS)

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# DNS protocol message format

**Additional**

'Glue' Resource Records (RRs) accompanying the authority section — so you can reach the name server without an extra lookup.

```
(a.ns.facebook.com, 129.134.30.12, A)
```

| identification                                                | flags                    | ← 12 bytes |
|---------------------------------------------------------------|--------------------------|------------|
| number of questions                                           | number of answer RRs     |            |
| number of authority RRs                                       | number of additional RRs |            |
| questions<br>variable number of questions                     |                          |            |
| answers<br>variable number of resource records                |                          |            |
| authority<br>variable number of resource records              |                          |            |
| additional information<br>variable number of resource records |                          |            |

# Query & response: a live example

## QUERY →

Transaction ID 0x6492  
Flags 0x0100 (query)  
Questions 1 Answer RRs 0  
Authority RRs 0 Additional RRs 0

### QUESTIONS

www.facebook.com A

## ← RESPONSE

Transaction ID 0x6492  
Flags 0x8180 (response, no error)  
Questions 1 Answer RRs 2  
Authority RRs 4 Additional RRs 8

### QUESTIONS

www.facebook.com A

### ANSWERS

www.facebook.com CNAME  
star-mini.c10r.facebook.com  
star-mini.c10r.facebook.com A  
31.13.78.35

### AUTHORITY

c10r.facebook.com NS  
d.ns.c10r.facebook.com (+3)

### ADDITIONAL

d.ns.c10r.facebook.com A  
185.89.219.11 (+7)

# DNS — key takeaways

- **Main purpose** — DNS maps human-readable hostnames to IP addresses.
- **Hierarchy** — DNS is distributed: no single server holds all records.  
Responsibility is delegated across Root, TLD, and Authoritative servers.
- **Traversal pattern** — the client queries the local DNS resolver, which then follows referrals Root → TLD → Authoritative.  
Results are cached locally by TTL to avoid repeated lookups.

## References

James F. Kurose & Keith W. Ross, *Computer Networking: A Top-Down Approach*, Addison Wesley, 3<sup>rd</sup> edition, 2005.

Chapter 2 Application Layer (2.1, 2.2, 2.4, 2.5)

William Stallings, *Data and Computer Communications*, Pearson, 10<sup>th</sup> edition, 2013.

Chapter 24 Electronic Mail, DNS, and HTTP (24.1–24.3)

Andrew S. Tanenbaum & David J. Wetherall, *Computer Networks*, Prentice Hall, 5<sup>th</sup> edition, 2010.

Chapter 7 The Application Layer (7.1–7.3)

# Acknowledgements

These lecture notes are developed based on slides and notes from:

Barry Wu's lecture notes for COSC264, University of Canterbury

Dr Mengmeng Ge's lecture notes for COSC264, University of Canterbury, 2024

Assoc Prof Dan Kim's lecture notes for COSC264, University of Canterbury, 2017

Prof Aleksandar Kuzmanovic's lecture notes for CS340, Northwestern University, 2005